
NATIONAL ECONOMICS UNIVERSITY (NEU)

Faculty of Mathematical Economics

RESEARCH REPORT — TIME SERIES ANALYSIS

Daily Temperature Forecasting

Using Time Series and Machine Learning Models

Author Nguyen Tran Tuan Kiet

Student ID 11230554

University National Economics University (NEU)

Date May 2026

Hanoi, Vietnam

Abstract

This paper analyzes and forecasts daily temperature in **Mumbai, India** using meteorological data recorded over the period 2016–2020 (1,781 observations). Three models are estimated and compared: a univariate SARIMA model, a multivariate SARIMAX with exogenous variables, and an XGBOOST gradient-boosting model. Exogenous variables for SARIMAX are selected through Granger causality testing, identifying wind direction, sea-level pressure, humidity, and dew point as statistically significant predictors. On the held-out test set (January–November 2020), SARIMAX achieves the highest forecast accuracy ($R^2 = 0.976$, RMSE = 0.313 °C, MAPE = 0.83%), outperforming both XGBOOST ($R^2 = 0.938$) and SARIMA ($R^2 = 0.836$). Results underscore the importance of causality-based variable selection for multivariate time series forecasting.

Keywords: time series forecasting, SARIMA, SARIMAX, XGBoost, Granger causality, Mumbai temperature prediction.

Contents

1	Introduction	3
2	Literature Review	3
2.1	Statistical Time Series Models	3
2.2	Granger Causality and Variable Selection	4
2.3	Machine Learning for Time Series Forecasting	4
3	Data and Methodology	4
3.1	Data Description	4
3.2	Preprocessing	6
3.3	Feature Engineering	8
3.4	Stationarity Testing	8
3.5	Granger Causality Testing	9
3.6	Model Specification	10
3.6.1	SARIMA	10
3.6.2	SARIMAX	10
3.6.3	XGBoost	11
3.7	Evaluation Framework	11
4	Results	11
4.1	SARIMA Results	11
4.2	SARIMAX Results	12
4.3	XGBoost Results	13
4.4	Model Comparison	14
5	Discussion	15
5.1	Granger Causality vs Correlation-Based Selection	15
5.2	SARIMAX vs XGBoost	16
5.3	Limitations	16
6	Conclusion	16

1. Introduction

Accurate temperature forecasting is essential for a wide range of economic and social applications, including agricultural planning, energy demand management, disaster preparedness, and public health. Mumbai, India — a tropical coastal city characterized by a distinct monsoon season (June–September) and strong annual temperature cycles — provides a particularly suitable case study for seasonal time series forecasting. Despite rapid advances in numerical weather prediction, statistical and machine learning approaches remain valuable for short-range forecasting at the local level, particularly when historical observation data are available but computational resources for physical models are limited.

This paper addresses the question: *which combination of model specification and variable selection best forecasts daily temperature using historical meteorological records?* The analysis makes three contributions. First, it applies Granger causality testing to identify exogenous predictors in a principled, data-driven manner, avoiding the common practice of selecting variables based on pairwise correlation alone. Second, it compares two statistical time series models (SARIMA and SARIMAX) with a widely used machine learning approach (XGBOOST), providing a comprehensive benchmark. Third, it uses a rolling one-step-ahead forecast methodology for SARIMA and SARIMAX to produce realistic out-of-sample forecasts, rather than the flat long-horizon predictions that arise from static multi-step forecasting.

The empirical results show that SARIMAX with Granger-selected variables substantially outperforms the univariate SARIMA baseline (RMSE reduction of 61.5%) and also outperforms XGBOOST, despite relying on fewer inputs and no non-linear interaction terms. This finding has practical implications for forecast model selection in meteorological and economic time series applications.

The remainder of the paper is organized as follows. Section 2 reviews related literature. Section 3 describes the data and methodology. Section 4 reports empirical results. Section 5 discusses findings and limitations, and Section 6 concludes.

2. Literature Review

2.1. Statistical Time Series Models

Time series forecasting of meteorological variables has a long tradition in statistics and atmospheric science. The Box–Jenkins framework (Box et al., 2015) provides a systematic approach to modeling univariate time series through ARIMA models. The seasonal extension, SARIMA, captures periodic patterns through multiplicative seasonal differencing and seasonal AR/MA components, making it well suited to variables such as temperature that exhibit strong annual cycles (Chatfield, 2000).

Multivariate extensions, including SARIMAX, incorporate exogenous variables to improve forecast accuracy when external predictors carry information beyond the target series’ own history (Hamilton, 1994). The performance of these models is sensitive to variable selection: including irrelevant predictors introduces multicollinearity and may worsen accuracy relative to the univariate baseline.

2.2. Granger Causality and Variable Selection

Granger (1969) introduced the concept of Granger causality as a formal test of predictive precedence: variable X is said to Granger-cause Y if past values of X improve the forecast of Y beyond what past values of Y alone can achieve. Formally, X Granger-causes Y if:

$$\sigma^2[\hat{Y}_{t+1} | \mathcal{F}_t(X, Y)] < \sigma^2[\hat{Y}_{t+1} | \mathcal{F}_t(Y)] \quad (1)$$

This criterion provides a principled basis for exogenous variable selection in SARIMAX, distinguishing genuine predictive relationships from spurious correlations driven by shared seasonal patterns. Granger causality tests have been widely used in macroeconometrics (Hamilton, 1994) but their application to meteorological variable selection for SARIMAX forecasting remains limited in the literature.

2.3. Machine Learning for Time Series Forecasting

Gradient-boosting algorithms, particularly XGBOOST (Chen and Guestrin, 2016), have demonstrated strong predictive performance across tabular forecasting tasks. For time series applications, these models are typically augmented with lag features and rolling statistics. XGBOOST’s ability to capture non-linear interactions makes it a natural benchmark for comparison with statistical models. Comparative studies generally find that machine learning models improve relative to statistical approaches as the number of observations and features grows, while statistical models retain advantages in low-data settings and for interpretability (Chatfield, 2000).

3. Data and Methodology

3.1. Data Description

The dataset consists of 1,781 daily meteorological observations recorded in Mumbai, India between 1 January 2016 and 15 November 2020. The target variable is `temp`, the daily mean temperature in degrees Celsius. Seven candidate predictor variables are available, as described in Table 1. Descriptive statistics are reported in Table 2.

Table 1: Data Dictionary

Variable	Type	Unit	Description
temp	float	°C	Daily mean temperature (target)
dew	float	°C	Dew point temperature
humidity	float	%	Relative humidity
sealevelpressure	float	hPa	Sea-level atmospheric pressure
winddir	float	°	Wind direction
solarradiation	float	—	Solar radiation
windspeed	float	km/h	Wind speed
preciptype	binary	0/1	Precipitation occurrence

Table 2: Descriptive Statistics (2016–2020, $n = 1,781$)

Variable	Mean	Std	Min	Median	Max
temp (°C)	28.34	1.96	20.20	28.50	32.80
dew (°C)	21.67	3.79	5.00	22.40	27.50
humidity (%)	69.58	15.22	28.60	70.30	98.20
sealevelpressure (hPa)	1008.91	3.12	994.10	1009.10	1017.40
winddir (°)	202.44	57.85	65.60	207.10	316.00
solarradiation	230.14	62.41	52.80	232.50	330.90
windspeed (km/h)	22.13	6.89	9.40	21.30	128.10
preciptype (0/1)	0.456	0.498	0	0	1

Temperature exhibits clear seasonal variation, with monthly means ranging from 25.4 °C in January to 30.8 °C in May (range: 5.4 °C). The distribution is slightly left-skewed ($\gamma_1 = -0.51$) with near-normal kurtosis ($\gamma_2 = 0.16$). The linear trend is negligible ($+0.067$ °C/year, $R^2 = 0.002$). Figure 1 shows the distribution and Figure 2 the full time series.




Figure 1: Distribution of daily temperature. Panel A: histogram with KDE overlay. Panel B: boxplot (23 outliers below lower IQR fence). Panel C: Q–Q plot against normal distribution.


Figure 2 — Temperature Time Series Analysis (2016–2020)


Figure 2: Temperature time series (2016–2020). Panel A: daily values with 30-day and 365-day moving averages and linear trend ($+0.067\text{ }^{\circ}\text{C}/\text{yr}$). Panel B: 30-day rolling standard deviation. Panel C: daily deviation from global mean ($28.3\text{ }^{\circ}\text{C}$).

3.2. Preprocessing

The dataset contains no duplicate rows, no duplicate timestamps, and no missing values. Outlier detection via the IQR method flags extreme values in several variables (Table 3). Winsorization at the 1st–99th percentiles is applied to `solarrradiation` (36 values clipped to $[84.74, 323.78]$), `sealevelpressure` (35 values to $[1000.40, 1015.72]$), and `dew` (30 values to $[9.60, 26.60]$). Temperature outliers (23 cold-season readings in January–February) and wind speed outliers (33 monsoon-season values) are retained as genuine meteorological events.

Table 3: Outlier Detection (IQR Method)

Variable	Lower fence	Upper fence	N outliers	%
temp	23.45	33.45	23	1.29
dew	7.30	36.30	7	0.39
sealevelpressure	999.80	1018.20	3	0.17
solarradiation	81.75	381.35	8	0.45
windspeed	5.55	37.55	33	1.85
humidity, winddir	—	—	0	0.00

**Figure 3:** Boxplots before and after outlier treatment. Winsorization eliminates outliers in solarradiation, sealevelpressure, and dew; temp and windspeed are retained.

3.3. Feature Engineering

Features are constructed in four groups targeting specific model classes (Table 4). StandardScaler normalization is applied after the train/test split, fitted on the training set only, to prevent data leakage. After dropping rows with missing values from lag operations (7 rows), the final dataset contains 1,774 observations and 22 features.

Table 4: Feature Engineering Summary by Model

Group	Features	SARIMAX	XGBoost
Time features	month, quarter, dayofyear, season, sin / cos(doy), sin / cos(month)	×	✓
Exogenous lags	dew_lag1, sealevelpressure_lag1	✓	✓
Target lags	temp_lag1 ($r = 0.908$), temp_lag7 ($r = 0.630$)	×	✓
Rolling mean	temp_rollmean7 ($r = 0.801$)	×	✓
Total		2	20

SARIMA uses no external features.




Figure 4: Lag and rolling features vs daily temperature. `temp_lag1` (red) tracks the series with a one-day offset; `temp_rollmean7` (green) provides a smoothed local trend.

3.4. Stationarity Testing

The Augmented Dickey–Fuller (ADF) test (Dickey and Fuller, 1979) is applied to the temperature series. Results are reported in Table 5.

Table 5: ADF Test Results for Daily Temperature

Series	ADF stat.	p -value	CV 1%	CV 5%	Decision
Original (y_t)	−4.617	0.0001	−3.434	−2.863	Stationary
First diff. (Δy_t)	−16.391	< 0.001	−3.434	−2.863	Stationary

The original series is stationary at the 1% level, rejecting the unit root null. The ACF and PACF of the first-differenced series (Figure 5) show significant spikes at lags 1–3 in both

functions (ACF lags 1–5: $[0.056, -0.130, -0.193, -0.087, -0.028]$; 95% CI: ± 0.047), suggesting candidate orders $p, q \in \{1, 2, 3\}$.




Figure 5: ACF and PACF of temperature after first differencing ($d = 1$). Red bars exceed the 95% confidence interval (± 0.047). Significant spikes at lags 1–3 guide the SARIMA grid search.

3.5. Granger Causality Testing

All seven candidate exogenous variables are subjected to pairwise Granger causality tests with respect to temperature, using up to 7 lags on first-differenced series. A VAR model is first estimated to guide lag selection; AIC selects an optimal lag order of 4, validating the choice of 7 as a conservative upper bound. Results are in Table 6 and Figure 6.

Table 6: Granger Causality Test Results (H_0 : Variable does not Granger-cause temp)

Variable	Best lag	Min p -value	Sig.	Decision
winddir	1	< 0.001	***	Granger-causes temp
humidity	3	0.0063	**	Granger-causes temp
sealevelpressure	1	0.0071	**	Granger-causes temp
dew	1	0.0315	*	Granger-causes temp
solarradiation	5	0.5802	ns	No effect
windspeed	3	0.2256	ns	No effect
preciptype	6	0.0720	ns	No effect

* $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$; ns = not significant

Notably, `winddir` is the most significant Granger-cause ($p < 0.001$) despite a modest Pearson correlation with temperature ($r = 0.23$), illustrating the value of causality-based over correlation-based variable selection.




Figure 6: Granger causality p -values by lag. Solid lines: significant at $\alpha = 0.05$. Dashed lines: non-significant. Red line: $\alpha = 0.05$ threshold.

3.6. Model Specification

3.6.1. SARIMA

The $\text{SARIMA}(p, d, q)(P, D, Q)_s$ model is:

$$\Phi_P(B^s) \phi_p(B) \Delta^d \Delta_s^D y_t = \Theta_Q(B^s) \theta_q(B) \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \sigma^2) \quad (2)$$

with $s = 12$, $d = D = 1$. Orders are selected by AIC-minimizing grid search over $p, q \in \{0, 1, 2\}$ and $P, Q \in \{0, 1\}$ (36 combinations). The optimal specification is $\text{SARIMA}(2, 1, 2)(0, 1, 1)_{12}$, with AIC = 3470.46 and BIC = 3502.03 (Table 7).

Table 7: Top 5 SARIMA Models by AIC

p	d	q	P	D	Q	s	AIC	BIC
2	1	2	0	1	1	12	3470.46	3502.03
2	1	2	1	1	1	12	3472.46	3509.30
2	1	1	0	1	1	12	3474.44	3500.76
2	1	1	1	1	1	12	3476.44	3508.02
1	1	2	0	1	1	12	3477.64	3503.96

Bold = selected model.

3.6.2. SARIMAX

SARIMAX extends $\text{SARIMA}(2, 1, 2)(0, 1, 1)_{12}$ with Granger-selected exogenous variables:

$$\Phi_P(B^s) \phi_p(B) \Delta^d \Delta_s^D y_t = \beta^\top \mathbf{x}_t + \Theta_Q(B^s) \theta_q(B) \varepsilon_t \quad (3)$$

where $\mathbf{x}_t = (\text{winddir}_t, \text{humidity}_t, \text{sealevelpressure}_t, \text{dew}_t)^\top$, standardized using training-set parameters. The fitted model achieves AIC = 789.87 and BIC = 842.49,

substantially lower than SARIMA.

3.6.3. XGBoost

XGBOOST (Chen and Guestrin, 2016) minimizes the regularized objective:

$$\mathcal{L} = \sum_i \ell(y_i, \hat{y}_i) + \sum_k \Omega(f_k), \quad \Omega(f) = \gamma T + \frac{1}{2} \lambda \|\mathbf{w}\|^2 \quad (4)$$

Hyperparameters: max_depth=3, learning_rate=0.05, subsample=0.7, reg_alpha=0.1, reg_lambda=1.5, with early stopping at 30 rounds. The model converges at iteration 298.

3.7. Evaluation Framework

Data are split chronologically: training set (8 Jan 2016 – 31 Dec 2019; $n = 1,454$; 82%) and test set (1 Jan 2020 – 15 Nov 2020; $n = 320$; 18%) (Figure 7). SARIMA and SARIMAX use rolling one-step-ahead forecasting; at each step the model is re-fitted on all history and produces $\hat{y}_t$ before appending actual y_t . Performance is assessed using MAE, RMSE, MAPE, and R^2 .




Figure 7: Chronological train/test split at 1 January 2020.

4. Results

4.1. SARIMA Results

The best-fitting SARIMA(2, 1, 2)(0, 1, 1)₁₂ model is estimated on 1,454 training observations. Residual diagnostics (Figure 8) confirm that residuals are broadly white noise, with the exception of minor seasonal autocorrelation at lags 12–13. The Ljung–Box test, however, rejects the null of no autocorrelation at lags 10 and 20 ($p < 0.001$), indicating residual structure that is not fully captured by the selected model order — a known limitation of SARIMA when seasonality is approximate. On the test set, SARIMA achieves MAE = 0.610 °C, RMSE = 0.813 °C, MAPE = 2.22%, and $R^2 = 0.836$ (Figure 9).




Figure 8: SARIMA residual diagnostics. Panel A: residuals over time. Panel B: residual distribution vs normal fit. Panel C: ACF of residuals (minor significant spikes at lags 12–13). Panel D: Q–Q plot.




Figure 9: SARIMA rolling forecast vs actual temperature (test set 2020). MAE = 0.610 °C, RMSE = 0.813 °C, $R^2 = 0.836$.

4.2. SARIMAX Results

Incorporating the four Granger-selected variables substantially improves forecast accuracy. Test MAE falls to 0.231 °C (−62.1%), RMSE to 0.313 °C (−61.5%), MAPE to 0.83% (−62.5%), and R^2 rises to 0.976 (Figure 10). The RMSE ratio of 0.32 indicates that the rolling forecast test error is well below the training error, consistent with the rolling methodology and the absence of overfitting. Convergence warnings arose at a subset of rolling steps but did not materially affect forecast quality.




Figure 10: SARIMAX rolling forecast vs actual temperature (test set 2020). MAE = 0.231 °C, RMSE = 0.313 °C, $R^2 = 0.976$.

4.3. XGBoost Results

XGBOOST achieves test $R^2 = 0.938$, MAE = 0.353 °C, RMSE = 0.501 °C, and MAPE = 1.29% (Figure 11). The RMSE ratio of 1.42 indicates mild overfitting relative to the training set (train RMSE = 0.353 °C), within acceptable bounds for an ensemble model. Feature importance analysis (Figure 12) reveals that `temp_lag1` accounts for approximately 38% of total gain, followed by `season` ($\approx 25\%$) and `temp_rollmean7` ($\approx 13\%$).




Figure 11: XGBOOST forecast vs actual temperature (test set 2020). MAE = 0.353 °C, RMSE = 0.501 °C, $R^2 = 0.938$.




Figure 12: XGBOOST feature importance (gain). `temp_lag1`, `season`, and `temp_rollmean7` are the three most important features, collectively accounting for $\approx 76\%$ of total gain.

4.4. Model Comparison

Table 8 summarizes performance across all three models, and Figure 13 shows the three forecast series against actual temperature over the test period.

Table 8: Out-of-Sample Forecast Performance (Test Set: Jan–Nov 2020)

Metric	SARIMA	SARIMAX	XGBOOST
Type	Statistical	Statistical	Machine Learning
Exogenous vars	None	4 (Granger)	20 features
AIC	3470.46	789.87	N/A
BIC	3502.03	842.49	N/A
Train RMSE (°C)	1.2212	0.9819	0.3533
Test MAE (°C)	0.6100	0.2310	0.3532
Test RMSE (°C)	0.8131	0.3129	0.5012
Test MAPE (%)	2.2155	0.8316	1.2899
Test R^2	0.8362	0.9758	0.9378
RMSE ratio	0.67	0.32	1.42
Rank	3rd	1st	2nd

RMSE ratio = Test RMSE / Train RMSE. Bold = best in column.


Figure — SARIMA vs SARIMAX vs XGBoost: Forecast Comparison (Test Set 2020)

Figure 13: Forecast comparison on the test set (January–November 2020). Each panel shows actual temperature (black) against the model forecast (dashed). SARIMAX tracks the observed series most closely across all seasons.

5. Discussion

5.1. Granger Causality vs Correlation-Based Selection

The most important methodological finding is that Granger causality testing identifies a different and more useful set of predictors than Pearson correlation. Wind direction has only a modest correlation with temperature ($r = 0.23$) but is the most significant Granger-cause ($p < 0.001$, lag 1), suggesting its predictive value operates through a dynamic lagged channel rather than contemporaneous co-movement. Conversely, solar radiation is moderately

correlated with temperature ($r = 0.22$) but fails the Granger test ($p = 0.58$), implying its association is driven by shared seasonal patterns already captured by the ARIMA structure. An earlier exploratory analysis using only `dew` and `sealevelpressure` (selected by correlation) yielded negligible improvement over SARIMA, confirming that principled variable selection is critical.

5.2. SARIMAX vs XGBoost

SARIMAX outperforms XGBOOST despite using fewer inputs (4 vs 20 features) and no non-linear interaction terms. This outcome suggests that for a target variable with strong autoregressive structure and well-defined seasonal periodicity, a correctly specified statistical model with causally-selected covariates can match or exceed a flexible machine learning approach. XGBOOST’s reliance on `temp_lag1` for 38% of its gain further suggests it largely approximates the same autoregressive structure as SARIMA through a different parameterization, without fully exploiting the multivariate information available to SARIMAX.

5.3. Limitations

Several limitations should be acknowledged. First, the rolling forecast strategy assumes that exogenous variable values are known at each forecast step; in genuine real-time applications these would need to be forecast independently. Second, SARIMAX optimization produced convergence warnings at a subset of rolling steps; future work should explore alternative optimizers or reduced variable sets. Third, the five-year sample limits the number of full seasonal cycles available for estimation. Fourth, the analysis does not account for potential structural breaks or climate non-stationarity over longer horizons.

6. Conclusion

This paper compares SARIMA, SARIMAX, and XGBOOST for daily temperature forecasting in Mumbai, India, using 1,781 daily meteorological observations from 2016 to 2020. The central contribution is the application of Granger causality testing to select exogenous variables for SARIMAX, identifying wind direction, sea-level pressure, humidity, and dew point as statistically significant predictors.

Incorporating these variables yields a 61.5% reduction in RMSE relative to univariate SARIMA, lifting R^2 from 0.836 to 0.976 and MAPE below 1%. SARIMAX also outperforms XGBOOST ($R^2 = 0.938$), which relies on 20 engineered features. These results demonstrate that principled variable selection via Granger causality is key to unlocking the full potential of multivariate time series models.

Future research could extend this framework by applying walk-forward cross-validation, testing longer forecast horizons, incorporating additional atmospheric predictors, and exploring hybrid SARIMAX–XGBOOST architectures combining statistical interpretability with

machine learning flexibility.

References

- Box, G. E. P., Jenkins, G. M., Reinsel, G. C., and Ljung, G. M. (2015). *Time Series Analysis: Forecasting and Control*. John Wiley & Sons, Hoboken, NJ, 5th edition.
- Chatfield, C. (2000). *Time-Series Forecasting*. Chapman & Hall/CRC, Boca Raton, FL.
- Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794.
- Dickey, D. A. and Fuller, W. A. (1979). Distribution of the estimators for autoregressive time series with a unit root. *Journal of the American Statistical Association*, 74(366):427–431.
- Granger, C. W. J. (1969). Investigating causal relations by econometric models and cross-spectral methods. *Econometrica*, 37(3):424–438.
- Hamilton, J. D. (1994). *Time Series Analysis*. Princeton University Press, Princeton, NJ.

Appendix: Additional Figures




Figure 14: Seasonal patterns. Panel A: monthly mean with ± 1 standard deviation band. Panel B: boxplots by month; April–June highlighted as peak-heat period (30–31 °C).




Figure 15: ACF of daily temperature (lags 1–60). ACF remains significant to lag 55 (ACF lag-1 = 0.91), reflecting strong temporal persistence.




Figure 16: Correlation analysis. Panel A: Pearson correlation matrix. Panel B: correlation of each variable with temperature (ranked). Dew point ($r = 0.52$) and sea-level pressure ($r = -0.32$) are the strongest correlates.




Figure 17: Scatter plots of top predictors vs temperature, colored by month. All correlations significant at $p < 0.001$.




Figure 18: Annual temperature statistics (2016–2020). Annual mean fluctuates narrowly between 28.08 °C (2016) and 28.68 °C (2018), consistent with a negligible warming trend.